

# Naveen Kumar

+91-6397930569 | [connect.naveenk@gmail.com](mailto:connect.naveenk@gmail.com) | [github.com/itsnaveenk](https://github.com/itsnaveenk) | [linkedin.com/in/itsnaveenk](https://linkedin.com/in/itsnaveenk)

## EDUCATION

### Institute of Engineering & Technology, Lucknow

Master of Computer Applications (MCA)

Jun 2024

CGPA: 8.19 / 10.0

## TECHNICAL SKILLS

**Languages:** Java, Python, SQL

**Frameworks & Concepts:** Spring Boot, FastAPI, Microservices, REST APIs, Spring Security, Hibernate, JPA

**Databases:** MySQL, MongoDB, Redis, Elasticsearch, Vector Databases

**Cloud & DevOps:** AWS (EC2, S3, CloudFront, SQS, Lambda), Azure, Docker, Apache Kafka, Git, Logstash, Kibana

**AI / Search:** RAG, OpenAI API, Embeddings, Semantic Search, Hybrid Search, GitHub Copilot, Claude Code, Cursor

## EXPERIENCE

### Adda247

Oct 2024 - Present

Software Development Engineer 1 | Gurugram, India

**TestPrime** - Subscription-based exam prep platform (Mock Tests, PYP, Quizzes, Current Affairs)

- Designed and implemented REST APIs for subscription management, order processing, and full membership lifecycle including coupon validation, payment verification, and cancellation.
- Implemented freemium content access control with paywall integration, dynamically locking premium content for free-tier users and triggering purchase flow.
- Integrated UPI AutoPay for recurring subscriptions, contributed to migrating 5M+ users from course-based to monthly subscription model.
- Designed exam personalization APIs allowing users to select target exams; content APIs dynamically filter results based on user preferences stored in MySQL.
- Integrated analytics event tracking (Meta, Google Analytics, MoEngage, Trackier) for key actions: login, paywall view, subscription purchase, and test lifecycle events and sync events in Google Cloud BigQuery.
- Revamped User Authentication Service with multi-tenancy, OTP mobile login, AES encryption, and 2-device session enforcement.

**Adda247 Platform** - Core backend infrastructure and product features

- Built Search Module using Elasticsearch across quizzes, articles, current affairs, job alerts, and magazines, processing 600K+ weekly queries with 40% lower search latency.
- Built and maintained RESTful microservices; integrated Apache Kafka to process millions of in-app events daily for real-time service communication.
- Reduced load times by 30% via Redis caching and improved query performance by 40% through MySQL & MongoDB schema indexing optimizations.
- Automated 700K+ record migration to Elasticsearch using Logstash ETL pipelines with near real-time sync.
- Cut manual operational effort by 80% through automation pipelines for license keys and internal reporting via Google Sheets API.
- Maintained CI/CD pipelines and used Kibana/Grafana for log analysis, achieving a 20% reduction in downtime through faster incident resolution.

## PROJECTS

### AI-Powered Product Search & Recommendation System

Spring Boot | Elasticsearch | OpenAI | Hybrid Search

- Built a hybrid search pipeline combining BM25 keyword search and KNN vector search in Elasticsearch, supporting 1.5M+ products synced monthly for a large-scale e-commerce catalog.
- Implemented LLM-driven query refinement to extract structured filters (category, price, attributes) and handle vague queries using multi-turn conversation context.
- Designed LLM-based reranking and grounded response generation with strict JSON schemas and hallucination safeguards to improve result relevance.
- Reduced embedding API costs by detecting unchanged products at ingestion time, skipping unnecessary re-embeddings.
- Implemented chat history summarization with S3 persistence to enable context-aware search while keeping token usage and latency low.